



**Department Of Computer Science And Engineering  
Indian Institute Of Technology Guwahati**

# **Optimizing Memory Performance In Multi-Core Systems Through Data Placement**

**Akshay Bhosale**  
*(234101006)*

Under the Guidance of  
**Dr. Aryabartta Sahu**

# Contents

- Introduction
- Challenges And Objectives
- Review Of Prior Works
- Proposed Approach
- Experiments And Results
- Conclusion And Future Work

# Introduction


- Efficient memory performance is essential for multi-core systems
- Modern applications place heavy demands on memory
- Important to understand how cores interact with memory
- Involves studying access behaviors, and the hierarchical structure of DRAM (Dynamic Random Access Memory)

# Introduction

## Address Mapping in DRAM

### Virtual to Physical Address Mapping

- When a program requests data it uses a virtual address
- This translation uses page tables



**Fig. 1 Virtual Address to Physical Address translation**

# Introduction

## Address Mapping in DRAM

### Physical Address To Hardware Address Mapping

Physical address is divided into components that specify the channel, rank, bank, row, and column

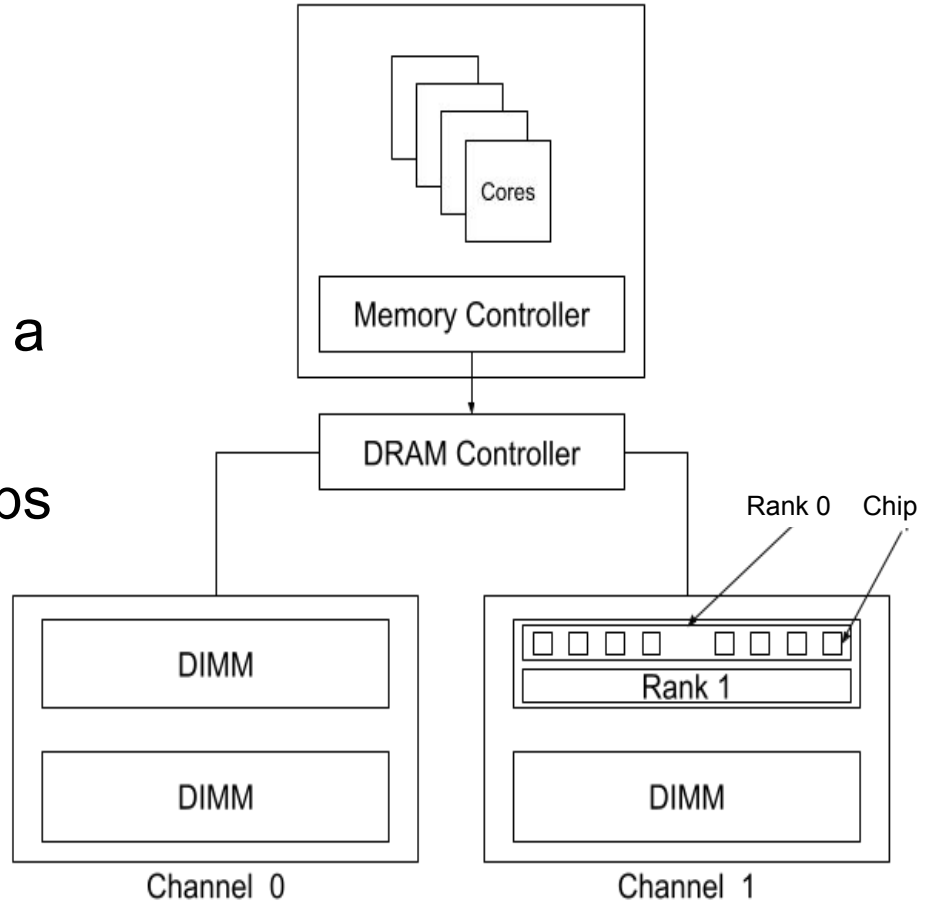


**Fig. 2 Physical Address breakdown**

# Introduction

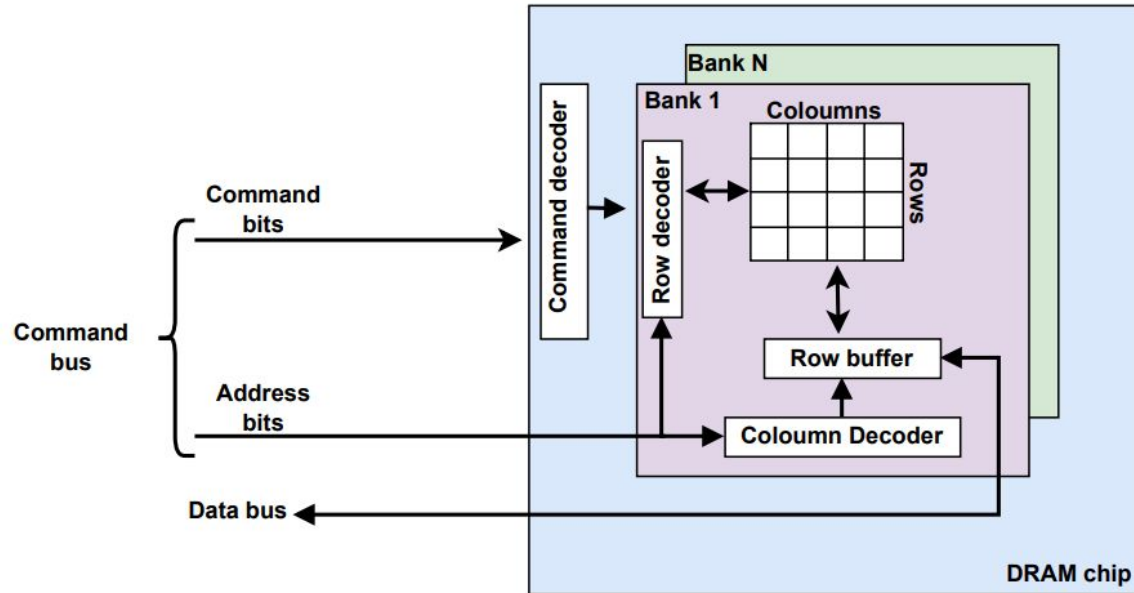
## DRAM Organization

- Organizes its components in a structured hierarchy
- Channels, DIMM, ranks, Chips banks, rows, and columns
- Each bank contains a row buffer



**Fig. 3 High-level architecture of DRAM**

# Introduction



**Fig. 4 DRAM Chip architecture**

# Introduction

## Shared Memory

- All processors share a common physical memory space
- Two types
  - UMA (Uniform Memory Access)
  - NUMA (Non-Uniform Memory Access)
- NUMA Improves scalability over UMA




# Introduction

## NUMA

- Memory is split across multiple memory channels
- Faster access to its local memory region
- Accessing remote memory is slower
- Latency varies depending on core-to-channel proximity

# Introduction

Uniform Memory Access (UMA)



Non Uniform Memory Access (NUMA)




Fig. 5 Differences between a NUMA and UMA

# Challenges

- Frequently accessed data may reside on high-latency channels
- Changing access patterns reduce placement effectiveness
- Overall number of remote accesses increases
- Results in suboptimal memory performance

# Objectives

- Reduce latency by minimizing remote accesses
- Place data closer to the cores that access it most frequently
- Adapt to changing access patterns in real-time
- Design a solution that is lightweight, modular, and practical

# Review Of Prior Work

- ML based memory and thread placement [1]
  - Uses profiling
- Replicate page table across nodes in NUMA [2]
  - targets metadata, not actual data
- Traffic-Aware Memory Placement in NUMA [3]
  - Reduce congestion
  - Collocation, replication and interleaving

# Review Of Prior Work

- Adaptive data placement in 3D-stacked DRAM [4]
  - Uses a small SRAM buffer
- Dynamically remaps physical address bits using access patterns [5]
  - Counters are used to monitor changes in bits
- Required hardware/kernel modifications or are tied to specific DRAM architectures

# Proposed Approach

- Map VPN to PPN in the memory channel closest to its dominant accessing core
- Periodically track and migrate hot pages to better channels
- Use Local to Requester and Dynamic Migration mapping
- Entirely at the translation layer no hardware changes needed

# Proposed Approach

## Local to Requester

---

### Algorithm 1 Local to Requester Mapping

---

**Require:** Virtual Page Number (VPN), Accessing Core ID (core\_id)

- 1: **if** VPN is already mapped **then**
  - 2:     **return** existing physical address
  - 3: **else**
  - 4:     channel  $\leftarrow$  closest memory channel to core\_id
  - 5:     candidate\_page  $\leftarrow$  select candidate page in channel
  - 6:     map VPN to candidate\_page
  - 7:     physical\_address  $\leftarrow$  combine candidate\_page with page offset
  - 8:     **return** physical\_address
  - 9: **end if**
-



# Proposed Approach

## Dynamic Migration

---

### Algorithm 2 Dynamic Migration

---

**Require:** Virtual Page Number (VPN), Accessing Core ID (core\_id)

- 1: Update per-core access count for VPN
  - 2: **if** migration window triggered **then**
  - 3:     Call migration module
  - 4: **end if**
  - 5: **if** VPN is already mapped **then**
  - 6:     **return** existing physical address
  - 7: **else**
  - 8:     Allocate physical page normally (e.g., Local to Requester)
  - 9:     Update mapping tables
  - 10:    **return** physical address
  - 11: **end if**
-

# Proposed Approach

## Migration module

---

```
1: for all hot pages do
2:   Determine dominant core and access distribution
3:   if a single dominant core exists then
4:     best_channel ← channel closest to dominant core
5:   else
6:     best_channel ← channel minimizing total latency across major cores
7:   end if
8:   if page is already in best_channel or cooldown is active then
9:     continue
10:  end if
11:  Calculate estimated gain and cost of migration
12:  if gain > cost then
13:    Migrate page to best_channel
14:    Update page table and cooldown timestamp
15:  end if
16: end for
```

---

# Experiment Setup

## System Configuration

- System with eight cores and eight DRAM channels.
- Each core is assigned to a nearby DRAM channel to capture NUMA memory access behavior
- DRAM type DDR4
- Simulated 1 billion instructions per benchmark to capture long-term behavior

# Experiment Setup




Fig. 6 Eight-core, eight-channel system used for simulation

# Experiment Setup

## Simulation Framework

- SniperSim performs simulations for multi-core and shared-memory workloads
- Ramulator2 is a detailed DRAM simulator modeling memory system behavior
- Benchmarks run on SniperSim to generate traces
- Traces are fed into Ramulator2 for simulation

# Experiment Setup

## Benchmarks

- Workloads from SPEC OMP2001
  - Art, Equake and FFT
- Other parallel programs
  - Gauss, PCNN and Matmul
- Covers diverse memory access patterns

# Experiment Setup

## Translation Policies Compared

- Random Mapping
  - Serves as a simple baseline for comparison
- Local to Requester
  - Captures the effect of initial placement near first accessing core
- Dynamic Migration
  - Tests the benefit of adaptive page migration
  - Nine migration configs by varying key parameters

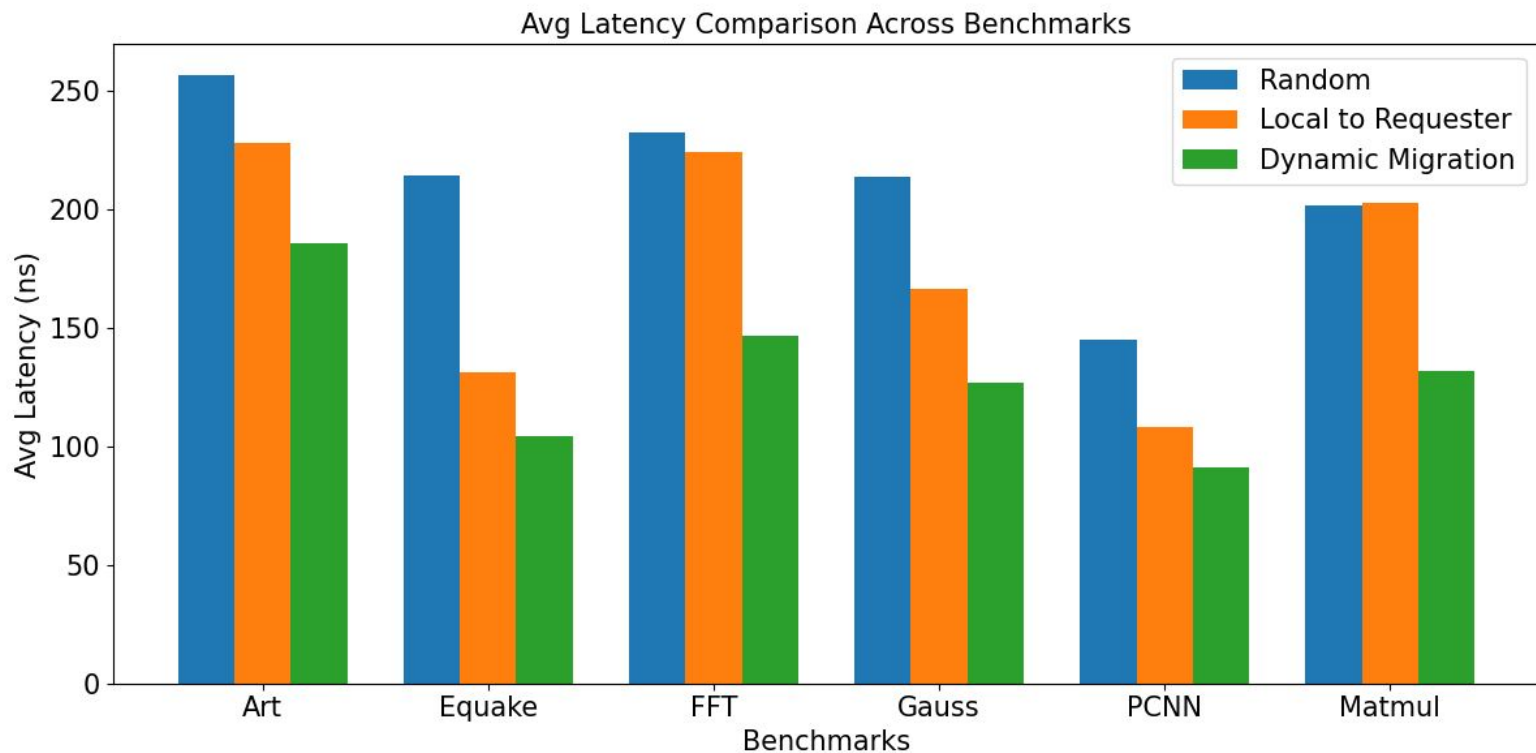
# Experiment Setup

## Dynamic Migration Configurations

- Parameters varied:
  - Window size
  - Hot page Threshold
  - Cooldown period
- Groups
  - Small windows: quick reaction
  - Large windows: more stable decisions



# Results and Analysis

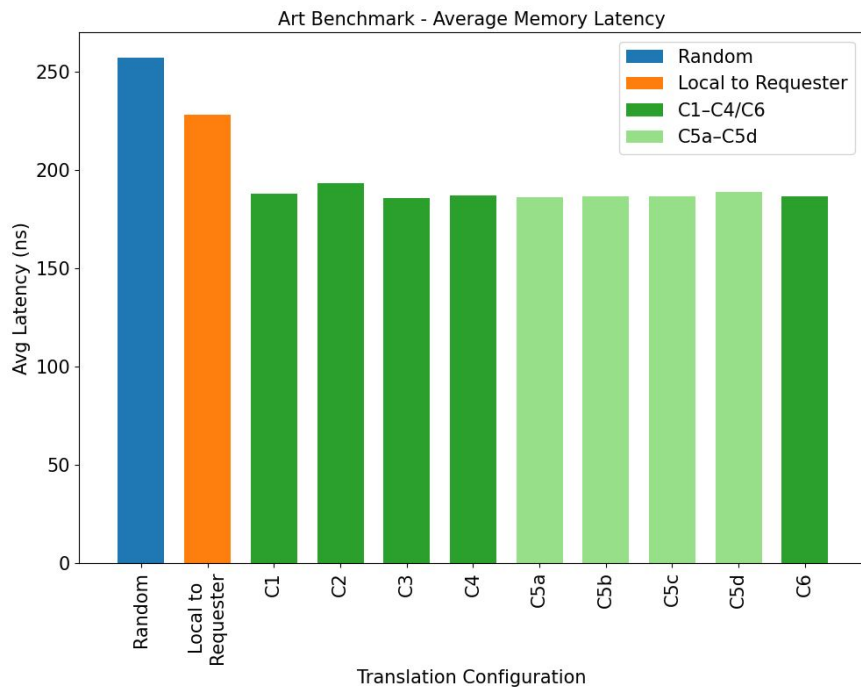


# Results and Analysis


## Translation Policies Compared

- Dynamic Migration consistently beats Random and Local to Requester
- Latency reduced by 27% to 50% across all benchmarks
- Biggest gains in Earthquake
- Shows benefit of adapting data placement at runtime

# Results and Analysis




Art Benchmark Avg Memory Latency

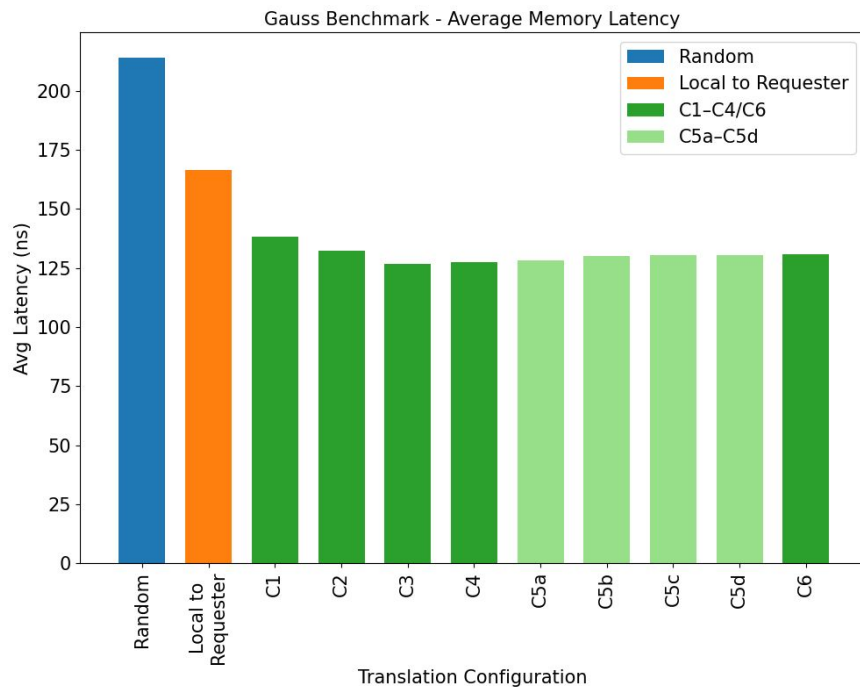


Equake Benchmark Avg Memory Latency

# Results and Analysis

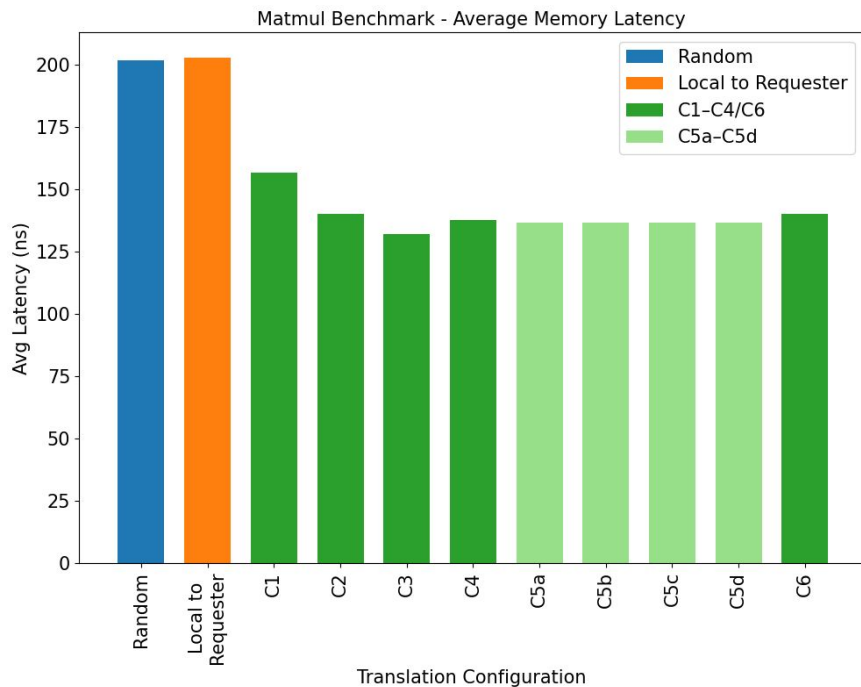


FFT Benchmark Avg Memory Latency




Gauss Benchmark Avg Memory Latency

# Results and Analysis



Matmul Benchmark Avg Memory Latency



PCNN Benchmark Avg Memory Latency

# Results and Analysis

## Sensitivity of Key Parameters

- All configs outperform Random and Local to Requester
- No config dominates
- Ensures meaningful decisions regardless of parameter tuning
- Fine-tuning yields minimal effect, system is stable and avoids overfitting

# Conclusion

- Lightweight runtime page placement at translation layer
- Dynamic Migration shows improvements ranging from 27% to over 50% in evaluated benchmark
- No hardware changes, practical and flexible
- Maintains performance regardless of different parameter settings

# Future Work

- Implement page replication with coherence support
- Integrate approach with real OS for practical deployment
- Explore energy optimization
- Explore integration with caching techniques for DRAM



# References

1. Denoyelle et al., “Data and Thread Placement in NUMA Architectures: A Statistical Learning Approach,” Proc. ICPP, 2019.
2. Panwar, “Efficient Virtual Memory and Page Table Management in Large NUMA and Virtualized Systems,” Ph.D. Thesis, IIT Kanpur, 2022.
3. Fedorova et al., “Carrefour: Traffic-Aware Memory Placement in Modern NUMA Systems,” ACM Queue, vol. 13, no. 7, 2015.
4. R. Pandey and A. Sahu, “Run-time adaptive data page mapping: A comparison with 3d-stacked dram cache,” Journal of Systems Architecture, vol. 110, p. 101798, 2020.
5. Qureshi et al., “DReAM: Dynamic Re-arrangement of Address Mapping to Improve Access Locality in DRAMs,” Proc. ISCA, 2010.
6. H. Luo, Y. C. Tuğrul, F. N. Bostancı, A. Olgun, A. G. Yağlıkçı, , and O. Mutlu, “Ramulator 2.0: A Modern, Modular, and Extensible DRAM Simulator,” 2023.
7. T. E. Carlson, W. Heirman, S. Eyerman, I. Hur, and L. Eeckhout, “An evaluation of high-level mechanistic core models,” ACM Trans. Archit. Code Optim., vol. 11, Aug. 2014.

# THANK YOU

**Submitted By:**

Akshay Bhosale (234101006)

**Under the Guidance of:**

Dr. Aryabartta Sahu (Associate Professor, CSE Dept, IIT Guwahati)

**Indian Institute of Technology, Guwahati**